

```
!pip install torch torchaudio torchvision transformers gradio Pillow

import gradio as gr
from transformers import pipeline, AutoModelForQuestionAnswering, AutoTokenizer
import torch
from PIL import Image
import json

# =====
# INPUT PROCESSOR - Handles multimodal inputs and speech-to-text conversion
# =====
class InputProcessor:
    def __init__(self):
        # Whisper for speech-to-text conversion
        self.stt = pipeline("automatic-speech-recognition", model="openai/whisper")

    def process_inputs(self, audio_file, text_input, image_file):
        """Process all input modalities and convert to standardized format"""
        processed_data = {
            'text_data': '',
            'image_data': None,
            'has_audio': False,
            'has_text': False,
            'has_image': False
        }

        # Process audio input - CONVERT AUDIO TO TEXT HERE
        if audio_file is not None:
            print(f"Processing audio file: {audio_file}")
            try:
                transcription = self.stt(audio_file)["text"]
                processed_data['text_data'] = transcription
                processed_data['has_audio'] = True
                print(f"Audio transcribed to: {transcription}")
            except Exception as e:
                print(f"Error transcribing audio: {e}")
                processed_data['text_data'] = f"Error processing audio: {str(e)}"

        # Process text input
        if text_input and text_input.strip():
            # Combine with audio transcription if both exist
            if processed_data['text_data']:
                processed_data['text_data'] += " " + text_input
                print(f"Combined audio + text: {processed_data['text_data']}")
            else:
                processed_data['text_data'] = text_input
                processed_data['has_text'] = True

        # Process image input
        if image_file is not None:
            processed_data['image_data'] = image_file
            processed_data['has_image'] = True
```

```

        print(f"Image processed: {image_file}")

    return processed_data

# =====
# GATING MECHANISM/ROUTER - Determines which expert models to engage
# =====
class GatingMechanism:
    def __init__(self):
        pass

    def route_to_experts(self, processed_data):
        """Determine which AI experts should be activated based on input data"""
        routing_plan = {
            'activate_clinical_llm': False,
            'activate_clinical_bert': False,
            'activate_llava_med': False,
            'activate_biomedclip': False,
            'data_for_experts': {}
        }

        # Route text data (including transcribed audio) to appropriate text-based
        if processed_data['text_data']:
            routing_plan['activate_clinical_llm'] = True
            routing_plan['activate_clinical_bert'] = True
            routing_plan['data_for_experts']['text'] = processed_data['text_data']
            print(f"Routing text to Clinical LLM and ClinicalBERT: {processed_data['text_data']}")

        # Route image data to vision experts
        if processed_data['has_image']:
            routing_plan['activate_llava_med'] = True
            routing_plan['activate_biomedclip'] = True
            routing_plan['data_for_experts']['image'] = processed_data['image_data']
            print("Routing image to LLaVA-Med and BioMedCLIP")

        return routing_plan

# =====
# AI HEALTHCARE EXPERTS - Specialized models for different modalities
# =====
class AIHealthcareExperts:
    def __init__(self):
        # Initialize ClinicalBERT for QA
        try:
            self.tokenizer = AutoTokenizer.from_pretrained("emilyalsentzer/BioMedCLIP")
            self.model_qa = AutoModelForQuestionAnswering.from_pretrained("emilyalsentzer/BioMedCLIP")
            self.clinical_bert = pipeline("question-answering", model=self.model_qa, tokenizer=self.tokenizer)
        except Exception as e:
            print(f"Error loading ClinicalBERT: {e}")
            self.clinical_bert = None

    def clinical_llm_analysis(self, text_data):

```

```

def clinical_llm_analysis(self, text_data):
    """Clinical LLM for general medical text analysis"""
    # Enhanced analysis with symptom extraction
    analysis = {
        'expert': 'Clinical_LLM',
        'findings': f"Clinical analysis of patient description: {text_data}"
        'confidence': 0.85,
        'key_symptoms': self._extract_symptoms(text_data)
    }
    print(f"Clinical LLM analysis completed with confidence: {analysis['confidence']}")
    return analysis

def _extract_symptoms(self, text):
    """Basic symptom extraction from text"""
    common_symptoms = ['pain', 'fever', 'cough', 'fatigue', 'headache', 'nausea']
    found_symptoms = []
    text_lower = text.lower()
    for symptom in common_symptoms:
        if symptom in text_lower:
            found_symptoms.append(symptom)
    return found_symptoms if found_symptoms else ['General discomfort mentioned']

def clinical_bert_qa(self, text_data):
    """ClinicalBERT for symptom extraction and medical QA"""
    if self.clinical_bert is None:
        return {
            'expert': 'ClinicalBERT',
            'qa_results': [{'question': 'Model unavailable', 'answer': 'ClinicalBERT is not available', 'confidence': 0.0}],
            'confidence': 0.0
        }

    # Generate relevant clinical questions
    questions = [
        "What symptoms is the patient experiencing?",
        "What is the patient's main complaint?",
        "Are there any concerning symptoms mentioned?",
        "What is the duration of symptoms?"
    ]

    results = []
    for question in questions:
        try:
            answer = self.clinical_bert(question=question, context=text_data)
            results.append({
                'question': question,
                'answer': answer['answer'],
                'confidence': answer['score']
            })
        except Exception as e:
            results.append({
                'question': question,
                'answer': f'Unable to extract: {str(e)}',
                'confidence': 0.0
            })

```

```

    avg_confidence = sum([r['confidence'] for r in results]) / len(results)
    print(f"ClinicalBERT QA completed with average confidence: {avg_confidence}")

    return {
        'expert': 'ClinicalBERT',
        'qa_results': results,
        'confidence': avg_confidence
    }

```

```

def llava_med_analysis(self, image_data):
    """LLaVA-Med for medical image analysis with natural language"""
    # Placeholder for LLaVA-Med - would use actual model
    analysis = {
        'expert': 'LLaVA_Med',
        'findings': 'Medical image analysis: Examining uploaded medical image',
        'confidence': 0.78,
        'detected_abnormalities': ['Image analysis placeholder - would detect abnormalities']
    }
    print(f"LLaVA-Med analysis completed with confidence: {analysis['confidence']}")
    return analysis

```

```

def biomedclip_classification(self, image_data):
    """BioMedCLIP for medical image classification"""
    # Placeholder for BioMedCLIP - would use actual model
    classification = {
        'expert': 'BioMedCLIP',
        'primary_diagnosis': 'Medical condition classification pending',
        'confidence': 0.82,
        'differential_diagnoses': ['Condition A', 'Condition B', 'Normal'],
        'probabilities': [0.45, 0.35, 0.20]
    }
    print(f"BioMedCLIP classification completed with confidence: {classification['confidence']}")
    return classification

```

```

# =====
# DIAGNOSTIC INTEGRATOR/SYNTHESIZER - Combines expert outputs
# =====

```

```

class DiagnosticIntegrator:
    def __init__(self):
        pass

    def synthesize_findings(self, expert_outputs):
        """Combine and synthesize findings from multiple AI experts"""
        # Extract symptoms from all expert outputs
        symptoms = []
        possible_conditions = []

        for output in expert_outputs:
            expert_name = output.get('expert', 'Unknown')

            # Extract symptoms from Clinical LLM
            if expert_name == 'Clinical_LLM':
                expert_symptoms = output.get('key_symptoms', [])

```

```

        symptoms.extend(expert_symptoms)

# Extract symptoms from ClinicalBERT Q&A
elif expert_name == 'ClinicalBERT':
    qa_results = output.get('qa_results', [])
    for qa in qa_results:
        if qa['confidence'] > 0.3 and 'symptom' in qa['question'].lower():
            answer = qa['answer'].lower()
            # Extract common symptoms from answers
            symptom_keywords = ['pain', 'fever', 'cough', 'fatigue']
            for keyword in symptom_keywords:
                if keyword in answer and keyword not in symptoms:
                    symptoms.append(keyword)

# Extract conditions from image analysis
elif expert_name in ['LLaVA_Med', 'BioMedCLIP']:
    diagnosis = output.get('primary_diagnosis', output.get('finding', ''))
    if diagnosis and 'placeholder' not in diagnosis.lower():
        possible_conditions.append(diagnosis)

# Remove duplicates and clean up
symptoms = list(set([s for s in symptoms if s and s != 'General discomfort']))

synthesis = {
    'symptoms': symptoms if symptoms else ['general discomfort'],
    'possible_conditions': possible_conditions,
    'medications': self._generate_medication_advice(symptoms),
    'home_care': self._generate_home_care_advice(symptoms),
    'when_to_see_doctor': self._generate_doctor_advice(symptoms)
}

print(f"Diagnostic integration completed: {len(symptoms)} symptoms identified")
return synthesis

def _generate_home_care_advice(self, symptoms):
    """Generate simple home care advice based on symptoms"""
    advice = []

    if 'cough' in symptoms:
        advice.extend(['drink warm liquids', 'use honey for throat'])
    if 'fever' in symptoms:
        advice.extend(['rest and stay hydrated', 'use cool compress'])
    if 'pain' in symptoms:
        advice.extend(['apply heat or cold as needed', 'gentle stretching'])
    if 'fatigue' in symptoms:
        advice.extend(['get plenty of sleep', 'eat nutritious meals'])

    # Default advice if no specific symptoms
    if not advice:
        advice = ['rest well', 'stay hydrated', 'eat healthy foods']

    return advice

def _generate_doctor_advice(self, symptoms):

```

```

    """Generate advice on when to see doctor"""
    urgent_symptoms = ['chest pain', 'shortness of breath', 'severe pain']

    for symptom in symptoms:
        for urgent in urgent_symptoms:
            if urgent in symptom:
                return ['see doctor immediately if symptoms worsen', 'call

    return ['see doctor if symptoms last more than a few days', 'contact dc

def _generate_medication_advice(self, symptoms):
    """Generate medication suggestions based on symptoms"""
    medications = []

    # Over-the-counter medications based on symptoms
    if 'cough' in symptoms:
        medications.append('cough drops or cough syrup')
    if 'fever' in symptoms:
        medications.append('paracetamol or ibuprofen for fever')
    if 'pain' in symptoms or 'headache' in symptoms:
        medications.append('paracetamol or ibuprofen for pain')
    if 'nausea' in symptoms:
        medications.append('anti-nausea medication if needed')
    if 'fatigue' in symptoms:
        medications.append('multivitamins to support energy')

    # For respiratory symptoms
    if 'shortness of breath' in ' '.join(symptoms) or 'breathing' in ' '.join(symptoms):
        medications.append('see doctor immediately - do not self-medicate')

    # For chest pain
    if 'chest pain' in ' '.join(symptoms):
        medications.append('seek immediate medical attention - do not self-

    # Default if no specific medications needed
    if not medications:
        medications = ['no specific medication needed - focus on rest and h

    return medications

# =====
# OUTPUT GENERATOR - Formats final diagnosis and recommendations
# =====
class OutputGenerator:
    def __init__(self):
        pass

    def generate_final_output(self, synthesis):
        """Generate simple, patient-friendly output"""
        output = {
            'symptoms': ' '.join(synthesis['symptoms']),
            'medications': ' '.join(synthesis['medications']),
            'home_care': ' '.join(synthesis['home_care']),
            'doctor_advice': ' '.join(synthesis['doctor_advice'])

```

```

        'doctor_advice': '', '.join(synthesis['when_to_see_doctor']),
        'formatted_report': self._format_simple_report(synthesis)
    }
    return output

def _format_simple_report(self, synthesis):
    """Format a simple, patient-friendly report"""
    report = "=== Your Health Summary ===\n\n"

    report += f"Symptoms you mentioned: {'', '.join(synthesis['symptoms'])}\n"

    report += "Medications you can try:\n"
    for med in synthesis['medications']:
        report += f"• {med}\n"
    report += "\n"

    report += "What you can do at home:\n"
    for advice in synthesis['home_care']:
        report += f"• {advice}\n"

    report += "When to see a doctor:\n"
    for advice in synthesis['when_to_see_doctor']:
        report += f"• {advice}\n"

    if synthesis['possible_conditions']:
        report += f"\nPossible conditions to discuss with doctor: {'', '.join(synthesis['possible_conditions'])}\n"

    report += "\n⚠ Important: This is AI advice only. Always talk to a real doctor.\n"

    return report

```

```

# =====
# MAIN ORCHESTRATION SYSTEM
# =====
class AIGPDoctorSystem:
    def __init__(self):
        self.input_processor = InputProcessor()
        self.gating_mechanism = GatingMechanism()
        self.ai_experts = AIHealthcareExperts()
        self.diagnostic_integrator = DiagnosticIntegrator()
        self.output_generator = OutputGenerator()

    def process_patient_case(self, audio_file, text_input, image_file):
        """Main processing pipeline following the architectural flow"""
        print("=== Starting AI GP Doctor Analysis ===")

        # Step 1: Input Processing (Audio → Text conversion happens here)
        processed_data = self.input_processor.process_inputs(audio_file, text_input, image_file)

        # Step 2: Gating/Routing
        routing_plan = self.gating_mechanism.route_to_experts(processed_data)

        # Step 3: Expert Analysis (Text from audio + manual text goes to models)

```

```

# Step 3: Expert Analysis (Route from data to relevant expert)
expert_outputs = []

if routing_plan['activate_clinical_llm']:
    output = self.ai_experts.clinical_llm_analysis(routing_plan['data_for_expert'],
    expert_outputs.append(output)

if routing_plan['activate_clinical_bert']:
    output = self.ai_experts.clinical_bert_qa(routing_plan['data_for_expert'],
    expert_outputs.append(output)

if routing_plan['activate_llava_med'] and 'image' in routing_plan['data_for_expert']:
    output = self.ai_experts.llava_med_analysis(routing_plan['data_for_expert'],
    expert_outputs.append(output)

if routing_plan['activate_biomedclip'] and 'image' in routing_plan['data_for_expert']:
    output = self.ai_experts.biomedclip_classification(routing_plan['data_for_expert'],
    expert_outputs.append(output)

# Step 4: Diagnostic Integration
synthesis = self.diagnostic_integrator.synthesize_findings(expert_outputs)

# Step 5: Output Generation
final_output = self.output_generator.generate_final_output(synthesis)

print("=== AI GP Doctor Analysis Complete ===")
return final_output

# =====
# GRADIO INTERFACE - FIXED: Removed 'optional' parameters
# =====
def create_gradio_interface():
    ai_system = AIGPDoctorSystem()

    def process_interface(audio_file, text_input, image_file):
        """Interface function for Gradio"""
        try:
            result = ai_system.process_patient_case(audio_file, text_input, image_file)

            # Return simple, patient-friendly outputs
            return (
                result['formatted_report'],
                result['symptoms'],
                result['medications'],
                result['doctor_advice']
            )
        except Exception as e:
            error_msg = f"Error processing request: {str(e)}"
            return (error_msg, error_msg, error_msg, error_msg)

    # Create Gradio interface - REMOVED 'optional=True' parameters
    demo = gr.Interface(
        fn=process_interface,
        inputs=[

```



```

        gr.Audio(type="filepath", label="Voice/Audio Input (will be converted to text)"),
        gr.Textbox(label="Textual Description", placeholder="Describe symptoms"),
        gr.Image(type="filepath", label="Medical Images")
    ],
    outputs=[
        gr.Textbox(label="Health Summary", lines=12),
        gr.Textbox(label="Your Symptoms", lines=2),
        gr.Textbox(label="Recommended Medications", lines=3),
        gr.Textbox(label="Doctor Advice", lines=2)
    ],
    title="AI GP Doctor - Simple Medical Advice",
    description="Tell me your symptoms (by voice or text) and I'll give you advice",
    examples=[
        [None, "I have chest pain and can't breathe well", None],
        [None, "I have a cough and feel tired", None],
    ]
)

```

```
return demo
```

```

# =====
# MAIN EXECUTION
# =====
if __name__ == "__main__":
    demo = create_gradio_interface()
    demo.launch()

```

```

/usr/local/lib/python3.11/dist-packages/huggingface_hub/utils/_auth.py:94:
The secret `HF_TOKEN` does not exist in your Colab secrets.
To authenticate with the Hugging Face Hub, create a token in your settings
You will be able to reuse this secret in all of your notebooks.
Please note that authentication is recommended but still optional to access
warnings.warn(

```

```

config.json: 100% 1.98k/1.98k [00:00<00:00, 35.7kB/s]

model.safetensors: 100% 290M/290M [00:03<00:00, 124MB/s]

generation_config.json: 100% 3.81k/3.81k [00:00<00:00, 73.0kB/s]

tokenizer_config.json: 100% 283k/283k [00:00<00:00, 3.45MB/s]

vocab.json: 100% 836k/836k [00:00<00:00, 6.52MB/s]

tokenizer.json: 100% 2.48M/2.48M [00:00<00:00, 7.61MB/s]

merges.txt: 100% 494k/494k [00:00<00:00, 3.81MB/s]

```

normalizer.json: 100%

52.7k/52.7k [00:00<00:00, 909kB/

s]

added_tokens.json: 100%

34.6k/34.6k [00:00<00:00, 378kB/

s]

special_tokens_map.json: 100%

2.19k/2.19k [00:00<00:00, 42.4kB/

s]

preprocessor_config.json: 100%

185k/185k [00:00<00:00, 2.88MB/

s]

Device set to use cpu

config.json: 100%

385/385 [00:00<00:00, 10.9kB/

s]

vocab.txt: 100%

213k/213k [00:00<00:00, 3.43MB/

s]

pytorch_model.bin: 100%

436M/436M [00:02<00:00, 166MB/

s]

model.safetensors: 100%

436M/436M [00:07<00:00, 48.2MB/

s]

Some weights of BertForQuestionAnswering were not initialized from the mode
You should probably TRAIN this model on a down-stream task to be able to us
Device set to use cpu

It looks like you are running Gradio on a hosted a Jupyter notebook. For th

Colab notebook detected. To show errors in colab notebook, set debug=True i
* Running on public URL: <https://e2dc7656b20d28ab52.gradio.live>

This share link expires in 1 week. For free permanent hosting and GPU upgra

AI Health Helper - Simple Medical Advice

Tell me your symptoms (by voice or text) and I'll give you simple advice on what you can do.

